

How Market Researchers Are Actually Using AI in 2026

Most market researchers in 2026 are using AI for three things: analyzing open-ended survey responses at scale, drafting discussion guides and screeners, and summarizing transcripts and findings. The hype has settled into something practical. Researchers aren't replacing their thinking with AI, but they are finding places where it saves real time.

I teach market researchers how to use AI, and I consult with teams trying to figure out where it fits. The pattern across industries is pretty consistent. Adoption is widespread but shallow. Most people are using ChatGPT or Claude for brainstorming and drafting. Far fewer have connected AI to their actual data or workflows in any systematic way.

Who's actually adopting AI in market research right now?

The split isn't about company size anymore. It's about whether someone on the team stepped up to figure it out.

Large companies in regulated industries (pharma, financial services, healthcare) are moving cautiously. They've got legal reviews, compliance checks, and procurement processes that slow everything down. Some have internal AI councils. Most are still in pilot mode.

Smaller agencies and independent consultants are moving faster on adoption but often without much governance. They're using free or pro accounts and figuring it out as they go. That's great for speed. It's risky for data handling, which I'll get into in the data security article.

The teams getting the most value have one thing in common: a champion. Someone who learned the tools, wrote the prompts, and brought everyone else along. Without that person, AI tools sit unused.

What specific tasks are researchers handing off to AI?

Three categories dominate right now.

First: open-end coding and theme extraction. This is the clear winner in 2026. Tools like Ascribe, Relative Insight, and even well-prompted general AI can process thousands of verbatim responses in minutes. The output still needs human review. But researchers consistently tell me it cuts coding time by more than half.

Second: survey and guide development. Researchers feed AI their objectives and get back first drafts of questionnaires, screeners, and discussion guides. Not finished products. But good enough to react to, which is faster than starting from scratch.

Third: summarization. AI pulls key themes from transcripts and writes first-pass summaries. Researchers treat this like an intern's draft. Useful as a starting point, but it needs a careful review before it goes anywhere near a client.

What's still not working?

Video and audio analysis tools promise a lot and deliver uneven results. Emotion detection, facial coding, tone analysis. The demos look impressive. The real-world performance is inconsistent enough that most researchers don't trust it yet.

Anything requiring deep cultural interpretation also falls short. AI can tell you what people said. It can't reliably tell you what they meant, especially across languages, regions, and cultural contexts. That judgment still sits with the researcher.

How is the researcher role actually changing?

The junior analyst role is changing fastest. Tasks that used to fill someone's first couple years (coding, transcript review, basic charting) are getting compressed. The upside is that junior researchers can start learning strategic thinking earlier. The risk is that we create researchers who never developed the instinct for reading raw data carefully.

Senior researchers aren't necessarily coasting through this. Many are struggling with something different: how to manage teams through the uncertainty. They're dealing with junior staff who are anxious about their roles, clients who want AI used everywhere (or nowhere), and leadership asking for an AI strategy they don't feel equipped to write. The technical skills are one thing. The people management around AI adoption is a whole separate challenge.

For more on how research roles are evolving, the Insights Association [publishes regular industry benchmarking reports](#) worth reading, and Greenbook has a section of [dedicated articles around changes in insights careers](#).

How should a research team get started if they haven't yet?

Start with one workflow, not a platform decision. Pick the most painful repetitive task on your team. For most, that's open-end coding or transcript summarization. Find a tool that does that one thing well. Use it for three months. Then decide what's next.

Don't start by buying an enterprise AI suite. The teams I've seen do that spend six months on configuration and governance before anyone actually uses anything. The teams that get value fastest start small, with tools researchers can try immediately, then build structure around what's working.

AI Training for Market Research Teams: What to Look For (and What to Skip)

The best AI training for market research teams teaches prompt engineering, output evaluation, and ethical data use, all through hands-on exercises using the actual tools and data types the team works with every day. Skip the vendor demos, the generic “AI for business” courses, and anything that doesn’t let participants build a skill they can use immediately.

I teach AI workshops for market researchers, and I’ve learned what works by watching what sticks and what doesn’t. Most AI training fails for the same few reasons. Here’s what to look for and what to avoid.

What does effective AI training for researchers actually cover?

Core skills. Not a curriculum with 47 modules.

First: prompt engineering in a research context. Not memorizing templates. Learning why a prompt works so you can adapt it when it doesn’t. This means teaching people to provide context, define the output format, specify what good looks like, and iterate. A prompt library is a starting point, not a skill.

Second: output evaluation. Researchers already know how to evaluate human-generated work. But AI mistakes look different. AI hallucinates citations, invents methodology details, and writes with confidence that’s inversely proportional to its accuracy. Training needs to show what those failure modes look like in real research outputs.

Third: for leaders especially, data governance. What goes into these tools and what doesn’t. This isn’t exciting, but it’s what keeps your legal team from shutting everything down. The [MRS Guidance on Using AI and Related Technologies](#) is a great reference for teams building this into their training.

What should teams skip?

Skip vendor demos dressed up as training. If more than 20 percent of the session is showing a specific tool’s interface, you’re being sold, not taught.

Skip anything that promises “AI certification.” There is no recognized certification for AI in market research. There are vendors who will give you a certificate with their logo on it. That’s marketing.

Skip training that doesn't use your industry's actual research scenarios. A generic workshop on "AI for business" won't help a pharma researcher who needs to write patient journey discussion guides. The exercises need to use the same kinds of briefs, data, and deliverables your team handles.

How long should the training take?

About 90 minutes per session, with hands-on exercises built in. Anything longer and people stop retaining information. Anything shorter and you can't get past the demo stage into actual skill-building.

The real learning happens after the session, in regular practice. The best approach is a series of focused 90-minute workshops, each on one skill, with time between them for people to practice and come back with questions. A single full-day session sounds efficient but rarely produces lasting behavior change. People walk out with notes they never look at again.

Who should deliver the training?

Someone who's done research, not someone who's done AI. The trainer needs to understand discussion guides, screeners, coding frames, and client deliverables. The AI knowledge is secondary.

If you're hiring externally, ask them to walk through a real research scenario during the selection process. Give them a mock brief and ask how they'd use AI at each stage. The ones who ask clarifying questions about the research before touching any tool are the ones you want.

What does success look like a month after training?

Your team should be using AI for at least one workflow without prompting from leadership. They should be able to explain what they do and don't trust AI to handle. And ideally, someone has become the unofficial AI lead who other people go to with questions.

If none of that is happening, the training didn't stick. That usually means it was too generic, too tool-focused, or wasn't followed by hands-on practice. Fix that before buying more.

What Data Should Market Researchers Never Put Into an AI Tool?

Market researchers should never put personally identifiable respondent data (names, locations, phone numbers, email addresses, IP addresses) into a public AI tool unless the tool has a contractual data processing agreement and you've confirmed it's active. If you're using a free or pro account on ChatGPT, Claude, or similar tools, assume everything you type is visible to the provider and could surface elsewhere.

This question comes up in almost every workshop I teach, especially from researchers in regulated industries. The answer changes as tools update their policies and enterprise features, but the principles stay consistent.

What counts as personally identifiable in a research context?

The obvious things: names, email addresses, phone numbers, physical locations, IP addresses. Strip these before anything goes near an AI tool.

But also watch for verbatim responses that include identifying details. A quote like "I've been a nurse at St. Mary's in Akron for 12 years" is potentially identifying, especially in smaller sample sizes. If you're uploading open-ended responses, scan for these kinds of details first.

Demographic combinations can also identify people indirectly. A 62-year-old female VP of marketing at a specific company in a specific city might be unique in your dataset. That's identifying information, even without a name attached.

What about business-sensitive information?

This is where freelancers and independent consultants need to be especially careful. If you're a solo researcher using a free or pro account on ChatGPT or Claude, don't put unpublished client findings, proprietary methodologies, or sensitive business strategy into those tools. Those accounts don't have data processing agreements. Your inputs could be used for training. That's a risk to your client relationships and potentially a breach of your contract. Enterprise accounts are different. They typically include agreements that prevent your data from being used for model training. But you have to be on the enterprise plan and confirm the settings are correct. Being on an enterprise plan with training data collection still enabled happens more often than you'd think.

Which tools are actually safe?

Enterprise versions of ChatGPT, Claude, and Google's Gemini offer data processing agreements. Research-specific tools like Quillit and Flowres have built-in protections designed for sensitive respondent data. Always look for a vendor's Privacy Policy, read it thoroughly, and look through their Trust Center documents if they have one.

The rule of thumb: if you haven't read the Privacy Policy and Terms and Conditions yourself, don't assume you're protected.

What about researchers in regulated industries?

Pharma, healthcare, financial services, and legal research teams face additional requirements. GDPR, HIPAA, and industry-specific regulations add layers of compliance that general AI tools may not address.

If you're in a regulated industry, your legal or compliance team almost certainly has guidance on AI tools. Find it. Read it. If it doesn't exist yet, ask for it. Researchers in these industries have told me that the biggest risk isn't using AI. It's using AI without checking what's already been approved or prohibited internally.

The [ICC/ESOMAR International Code](#) on market, opinion, and social research and data analytics is the definitive external reference here. It covers data minimization, privacy, transparency, and human oversight in the context of AI. The [MRS AI Guidance](#) is another resource worth bookmarking, especially for UK and European practitioners.

What's the practical checklist before uploading anything?

Four questions to ask yourself:

One: does this contain PII (name, email, phone, location, IP address)? If yes, strip it first or don't upload.

Two: does this contain unpublished client business information? If yes and you're on a free or pro account, don't upload.

Three: does this contain information that's sensitive under industry regulations? If yes, check with compliance before uploading.

Four: do I know whether this tool has a data processing agreement that covers my use case? If no, find out before uploading.

How to Write an AI Acceptable-Use Policy for a Market Research Team

A good AI acceptable-use policy for a market research team lists which tools are approved for which data types, specifies what researchers can and can't input, and explains how to flag problems. It should fit on two pages. The goal is to let your team use AI productively without exposing respondent data or client information to unnecessary risk.

I haven't written these policies for teams myself, but I've studied what makes them work and what makes researchers ignore them. The pattern is clear: short, specific, practical policies get followed. Long, vague, legalistic ones get skimmed and forgotten.

Why do most AI policies fail?

Because they're written by people who don't understand how researchers actually work.

The typical policy bans everything specific and allows nothing. "Don't put client data into AI tools." But researchers handle client data all day. They need to know what's actually prohibited and what's fine. A policy that's too broad gets ignored because following it would mean doing nothing at all.

The other failure mode is length. If your policy is 15 pages, nobody reads it. The most effective policies I've seen are short enough to discuss in a team meeting and specific enough that someone knows what to do on a Tuesday afternoon.

What are the essential sections?

Four sections cover almost everything.

Section one: approved tools. List them by name and tier. ChatGPT Enterprise, yes. Consumer ChatGPT, no. Claude Team plan, yes. Free Claude, no. Be specific. Update this quarterly because the landscape changes fast.

Section two: data classification. Three buckets work well. Green for data you can put anywhere (public reports, your own prompts, non-sensitive methodology notes). Yellow for data that needs an approved tool with a data processing agreement (anonymized aggregate findings, draft deliverables). Red for data that never goes into any AI tool (raw PII, client financials, unpublished strategic findings). Give concrete examples in each bucket.

Section three: the approval process. Who can add a new tool to the approved list? What's the review process? How long does it take? If this takes six months, researchers will find workarounds.

Section four: what to do when something goes wrong. Who do you tell? What information do they need? Make this feel like a normal process, not a disciplinary procedure.

Should the policy cover how AI output is used in deliverables?

Yes, absolutely. State that AI-generated content in client deliverables must be reviewed by a human and that researchers are responsible for accuracy regardless of the source. This is becoming a client expectation, and some RFPs now ask about it directly.

How often should the policy be updated?

Every six months, or whenever a major tool changes its terms. Assign one person to own the policy. That person should be someone who uses AI regularly and stays current on the tools. Give them the authority to make minor updates without running a full approval process.

What's a workable outline?

Here's what I'd recommend:

1. Purpose and scope (two sentences)
2. Approved tools list (table format)
3. Data classification guide (three buckets with examples)
4. Input rules by data type (bullet points)
5. Output rules (human review required, accuracy responsibility)
6. New tool approval process (steps and timeline)
7. Incident reporting (who, what, when)
8. Policy owner and review schedule

If you're on page three, you're probably overdoing it. The [MRS Guidance on AI and Related Technologies](#) and the [ICC/ESOMAR International Code](#) are both worth referencing as external standards when building your policy. The [Insights Association's Data Quality and Standards resources](#) also provide useful framing, particularly around data integrity.

Prompt Frameworks for Market Research: A Starter Guide

The most useful prompt frameworks for market researchers are role-based prompts, chain-of-thought prompts, and structured output prompts. Master these three and you can handle about 90 percent of what research teams actually use AI for. You don't need to learn 20 different frameworks with clever acronyms.

I teach these frameworks in my workshops, and they're the ones participants actually use afterward. Here's the practical version. No theory. Just the patterns that work.

What's a role-based prompt and when do I use it?

You tell the AI who to be, what context it has, and what you need.

Use this for role playing - seriously. The magic is in the context. "You're a chief marketing officer" is too vague. "You're a chief marketing officer of a global CPG company needing to decide the direction of the next global ad campaign" is much stronger for reviewing your report to find the gaps and weaknesses.

The format: role, context, task, format, constraints. Like this:

You are a chief marketing officer of a global technology company listening to a recommendation on the messaging for the next global ad campaign that will run as part of the launch of a new service. The goal of the campaign is to increase purchase likelihood of the new service. The meeting is only 30 minutes, and my presentation should only be 10-15 of those minutes, with the rest of the time left for discussion. Identify where in the report the language and recommendations could be tighter, and list questions I could expect the CMO to raise during the meeting.

What's chain-of-thought prompting?

You ask the AI to show its reasoning before giving the answer. This matters for analysis tasks where the logic matters as much as the conclusion.

Use this for segmentation analysis, competitive positioning, and *anything* where you need to audit how the AI got to its answer. The prompt structure is simple: present the data or scenario, then say "Think through this step by step" before asking your question.

For example:

Here are verbatim responses from 50 consumers about why they switched coffee brands. Think through this step by step: first identify the main themes, then group the responses by theme,

then identify which themes appear most frequently. After that, tell me the top three reasons for brand switching with example quotes for each.

This catches errors faster because you can see where the reasoning went wrong instead of just seeing a wrong conclusion.

What's a structured output prompt?

You define exactly what the output should look like. Format, length, section headers, everything.

Use this for reporting, slide outlines, and anything that needs to fit into an existing deliverable format. The more specific you are, the less editing you'll do afterward. In fact, if you have a template or an example of a previous output that fits the structure you need (and that doesn't have business sensitive information), add it into the context for the AI to follow.

Instead of "summarize these findings," try:

Create a three-slide summary of these research findings. Slide one: key numbers (3-5 statistics, one sentence of context each). Slide two: three strategic implications (one paragraph each). Slide three: two recommended actions with rationale.

What about the other frameworks?

Most of the acronym-based frameworks (RACE, CARE, and the rest) are repackaged versions of these three patterns. You don't need to memorize them. You need to understand what each element does.

How do I get better at this?

Keep a prompt journal.

Every time you write a prompt that works well, save it. Note what the task was, what the prompt was, and what made the output good. Every time a prompt fails, save that too and note what went wrong. After a month, you'll have a personal library that's more useful than any generic prompt collection because it's calibrated to your actual work.

Share prompts with your team. The teams that get the most out of AI are the ones where prompt-sharing is part of the culture. Ethan Mollick's work at Wharton on practical AI adoption and prompt craftsmanship is a useful external reference. His [One Useful Thing](#) blog is one of the best resources out there for understanding how to actually work with AI, not just read about it.

Where's a good starting set?

Here are four prompts you can adapt today:

For discussion guides: "I'm preparing for focus groups about [topic] with [audience]. The research objective is [objective]. Draft a 90-minute discussion guide with timing, key questions, and suggested probes."

For open-end coding: "Here are [number] open-ended responses to the question '[question]'. Group these into 5-8 themes. For each theme, provide a label, a one-sentence description, the count of responses in that theme, and two representative quotes. Quotes must be verbatim from what is provided; do not summarize or combine quotes."

For report outlines: "Based on these findings [paste summary], create an outline for a 15-slide research report. Include a recommended slide title and 2-3 bullet points for each slide. The audience is [audience] and they care most about [priority]."

For survey critique: "Here is a draft survey about [topic] to be fielded with [audience]. Review it for leading questions, double-barreled questions, missing answer options, and logical flow issues. List each problem with the question number and a suggested fix."